

LAWRENCE FENG

☎ 925-819-8690

✉ lawrence.feng2017@gmail.com

lawrencefeng17.github.io

Education

Carnegie Mellon University

Pittsburgh, Pennsylvania

M.S. in Machine Learning, Expected May 2027

Beginning August 2026

B.S. in Statistics and Machine Learning; Minor in Computer Science; GPA: 3.90

May 2026

- Selected coursework: Introduction to Deep Learning; Deep Reinforcement Learning; Probabilistic Graphical Models; Probability and Mathematical Statistics; Algorithm Design and Analysis; Advanced Methods for Data Analysis.

Publications

[Early Data Exposure Improves Robustness to Subsequent Fine-Tuning](#). Lawrence Feng, Gaurav R. Ghosal, Jacob Mitchell Springer, Ziqian Zhong, and Aditi Raghunathan. *Conference on Language Modeling (COLM)*, 2026.

Research Experience

Research Assistant

Summer 2026 – Present

Carnegie Mellon Machine Learning Department; advised by Professor Aditi Raghunathan

- Solo-authored the technical report “[Subliminal Learning Happens at Every Rank, Given the Right Learning Rate and Enough Data](#)”, identifying and addressing methodological issues in prior work.
- Found that higher LoRA ranks are less data-efficient than lower ranks in this setting; ongoing work investigates why.

Honors Undergraduate Research Thesis

Fall 2025 – Spring 2026

Advised by Professor Aditi Raghunathan

- Led the project and conducted all experiments for a first-author COLM 2026 paper.
- Reframed catastrophic forgetting as a property of upstream training rather than solely a downstream mitigation problem: *how* a model acquires a capability matters for how it is retained, not just whether it acquires it.
- Introduced *early exposure*—mixing target data into pretraining—and demonstrated a *latent* benefit: it mitigates forgetting during downstream fine-tuning even when it provides no measurable benefit before fine-tuning.
- Developed theory showing that early exposure protects specialized features from downstream interference.

Watermarking for Language Models

Fall 2024 – Summer 2025

Carnegie Mellon Computer Science Department; advised by Professor Wenting Zheng

- Developed a more efficient sampling algorithm for embedding provably quality-preserving watermarks into LLM outputs.
- Owned experimental design, implementation, analysis, and progress reporting as the project’s sole programmer.
- Synthesized results into clear visualizations and technical reports for collaboration with non-ML researchers.

Robotics Institute at Carnegie Mellon

Spring 2024 – Fall 2024

Summer Research Apprenticeship

- Explored adapting CoTracker’s retrospective dense point tracking into predictive models that output future robot states.
- Uncovered significant robustness failures in state-of-the-art models through systematic environmental variations.

Industry Experience

Inflection AI

Summer 2025

Member of Technical Staff Intern

- Implemented [Qwen2.5-VL \(PR #2868\)](#) within the open-source torchtune library, enabling distributed fine-tuning.
- Designed an automatic benchmark curation tool that increased the statistical power of LLM-as-a-judge evaluations of stylistic properties of LLM outputs; this tool was adopted into the company’s core evaluation suite.
- Contributed to a company-wide emotional-intelligence evaluation effort by surveying prior work, improving benchmarks, prototyping persona-based user simulators, and building synthetic-data pipelines.

Leadership

Carnegie AI Safety Initiative

Fall 2023 – Present

President (Fall 2025 – Spring 2026), Member

- Spearheaded outreach growing mailing list 4× to ~ 900, expanded leadership team 4× to ~ 16, developed and taught [technical AI safety fundamentals course](#) to ~ 40 ppl/sem, piloted and grew jailbreaking hackathon to ~ 200 participants

Technical Skills

Programming Languages: Python, R, C, C++, Java, SQL, SML, L^AT_EX, Bash

Developer Tools: GCP, AWS EC2, Git, Weights & Biases, Slurm, Conda, uv

Frameworks and Libraries: PyTorch, TensorFlow, Hugging Face, vLLM, OpenCV, NumPy, Pandas, Matplotlib, ROS